

Database Indexing

Recap

- ❑ Our goal is to efficiently find the records specified in the 'where' clause of a query
- ❑ Binary search is a fast search algorithm
- ❑ Applying binary search to a table stored on disk has some problems:
 1. we can only sort the records of a table according to one key
 2. even on that one key, how to do the binary search on disk blocks?

Idea: create separate index table

Think of an index at the end of a book:

- it is stored separately from the rest of the book
- the entries are ordered for fast search
- it takes up space
- entries refer to one or more page numbers

Index for a table

index for 'name' field

search key	records
Brandt	
Califieri	
Crick	
Einstein	
...	

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

- **search key** is an attribute used to look up records
- index is ordered, so fast binary search can be used

Using the index

SQL: `select salary from instructor where name="Crick";`
`CREATE INDEX I1 ON INSTRUCTOR(NAME);`

Perform binary search on index

Get records from file using record ptrs in index

instructor

index for 'name' field

search key	records
Brandt	
Califieri	
Crick	
Einstein	
...	

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

© 2010 Blackwell Publishing Ltd *Journal of Internal Medicine* 267: 103–111

Index maintenance

Question: What to do if record inserted?

Question: What to do if instructor record deleted?

index for 'name' field

search key	records
Brandt	
Califieri	
Crick	
Einstein	
...	

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Index pros/cons

Pros:

- allows for binary search on multiple fields in a table

Cons:

- space to store indexes
- time to update indexes as table changes

Example: index on a primary key

In this example, index has same number of rows as instructor table

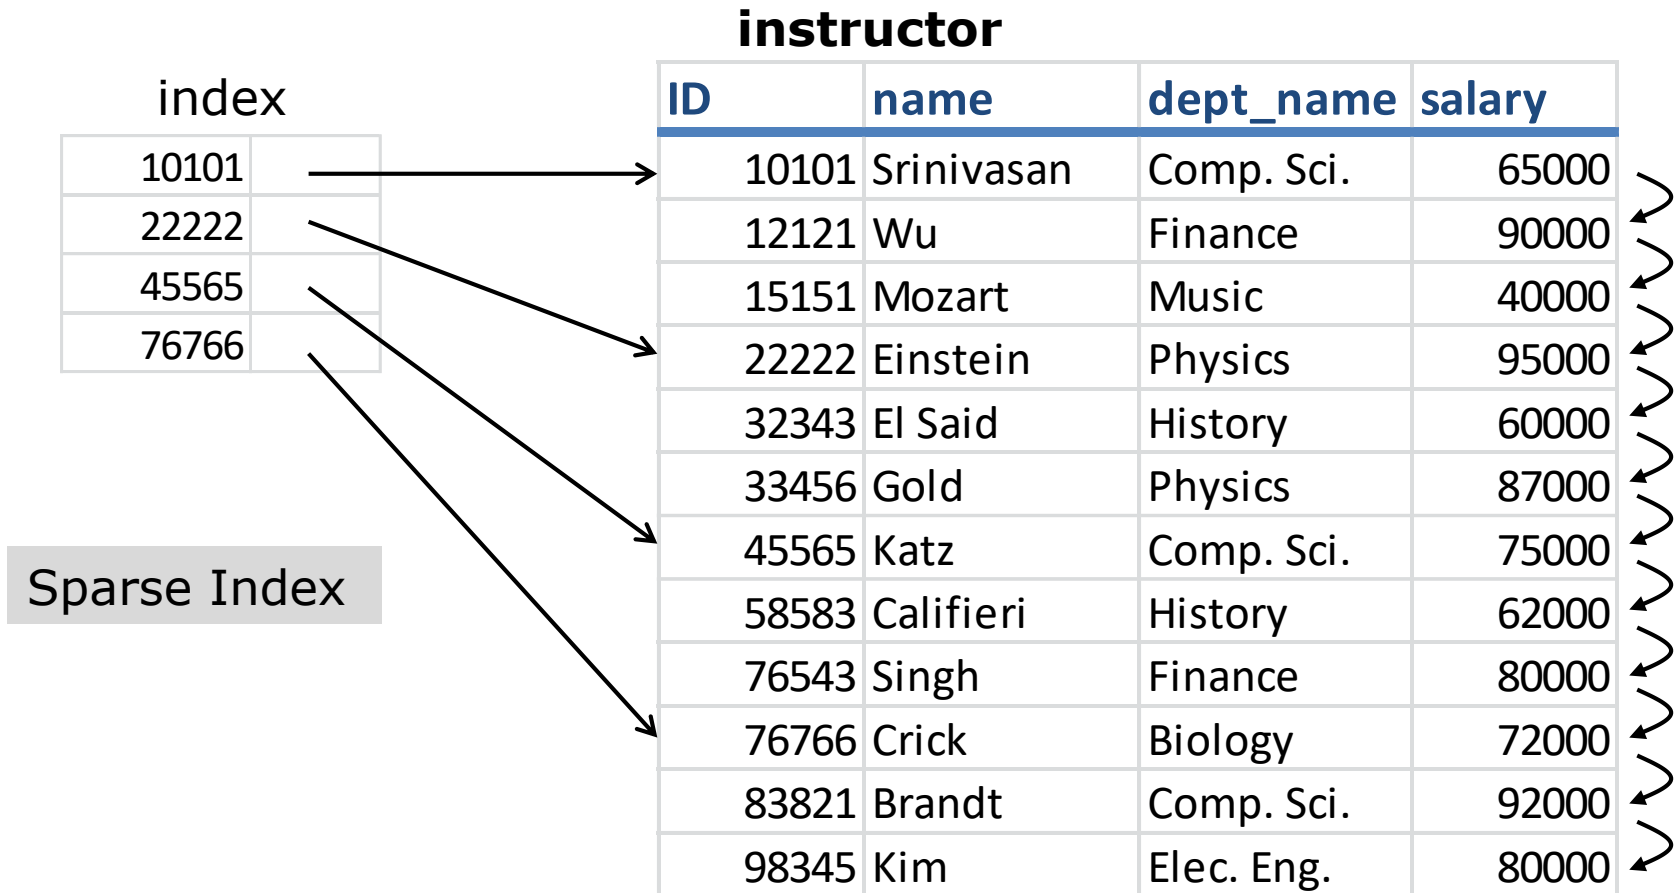
ID is the "search key"

Primary Index - Clustered

instructor

index		ID	name	dept_name	salary
10101	→	10101	Srinivasan	Comp. Sci.	65000
12121	→	12121	Wu	Finance	90000
15151	→	15151	Mozart	Music	40000
22222	→	22222	Einstein	Physics	95000
32343	→	32343	El Said	History	60000
33456	→	33456	Gold	Physics	87000
45565	→	45565	Katz	Comp. Sci.	75000
58583	→	58583	Califieri	History	62000
76543	→	76543	Singh	Finance	80000
76766	→	76766	Crick	Biology	72000
83821	→	83821	Brandt	Comp. Sci.	92000
98345	→	98345	Kim	Elec. Eng.	80000

Omitting some values from the index




Index on a different field

index		instructor			
		ID	name	dept_name	salary
Biology	→	76766	Crick	Biology	72000
Comp. Sci.	→	10101	Srinivasan	Comp. Sci.	65000
Elec. Eng.	↘	45565	Katz	Comp. Sci.	75000
Finance	↘	83821	Brandt	Comp. Sci.	92000
History	↘	98345	Kim	Elec. Eng.	80000
Music	↘	12121	Wu	Finance	90000
Physics	↘	76543	Singh	Finance	80000
		32343	El Said	History	60000
		58583	Califieri	History	62000
		15151	Mozart	Music	40000
		22222	Einstein	Physics	95000
		33456	Gold	Physics	87000

Note that the index has fewer rows than the table, but all search key values are in the index.

When are multiple pointers needed?



Recap

We've considered various methods for searching data on disk

- Linear search
- Binary search
- Binary search with ordered indexes

Ordered index pros and cons:

- + fast
- complexity
- space overhead
- time overhead when records added/deleted

Reminder: dense index

		instructor			
index		ID	name	dept_name	salary
10101	→	10101	Srinivasan	Comp. Sci.	65000
12121	→	12121	Wu	Finance	90000
15151	→	15151	Mozart	Music	40000
22222	→	22222	Einstein	Physics	95000
32343	→	32343	El Said	History	60000
33456	→	33456	Gold	Physics	87000
45565	→	45565	Katz	Comp. Sci.	75000
58583	→	58583	Califieri	History	62000
76543	→	76543	Singh	Finance	80000
76766	→	76766	Crick	Biology	72000
83821	→	83821	Brandt	Comp. Sci.	92000
98345	→	98345	Kim	Elec. Eng.	80000

In this example, index has same number of rows as instructor table


Search with dense index

Search process:

- given search key value
- perform binary search on index
- when index entry found, go to record in a block of table

If **index is small**, it can be kept in memory.

This search is fast.



What happens when table is large?


Suppose table has 1 million rows.

If 100 index entries per disk block, then index is 10,000 blocks.


Index is probably too big for memory.

Binary search on index now requires about 14 ($\log_2 10,000$) disk accesses.

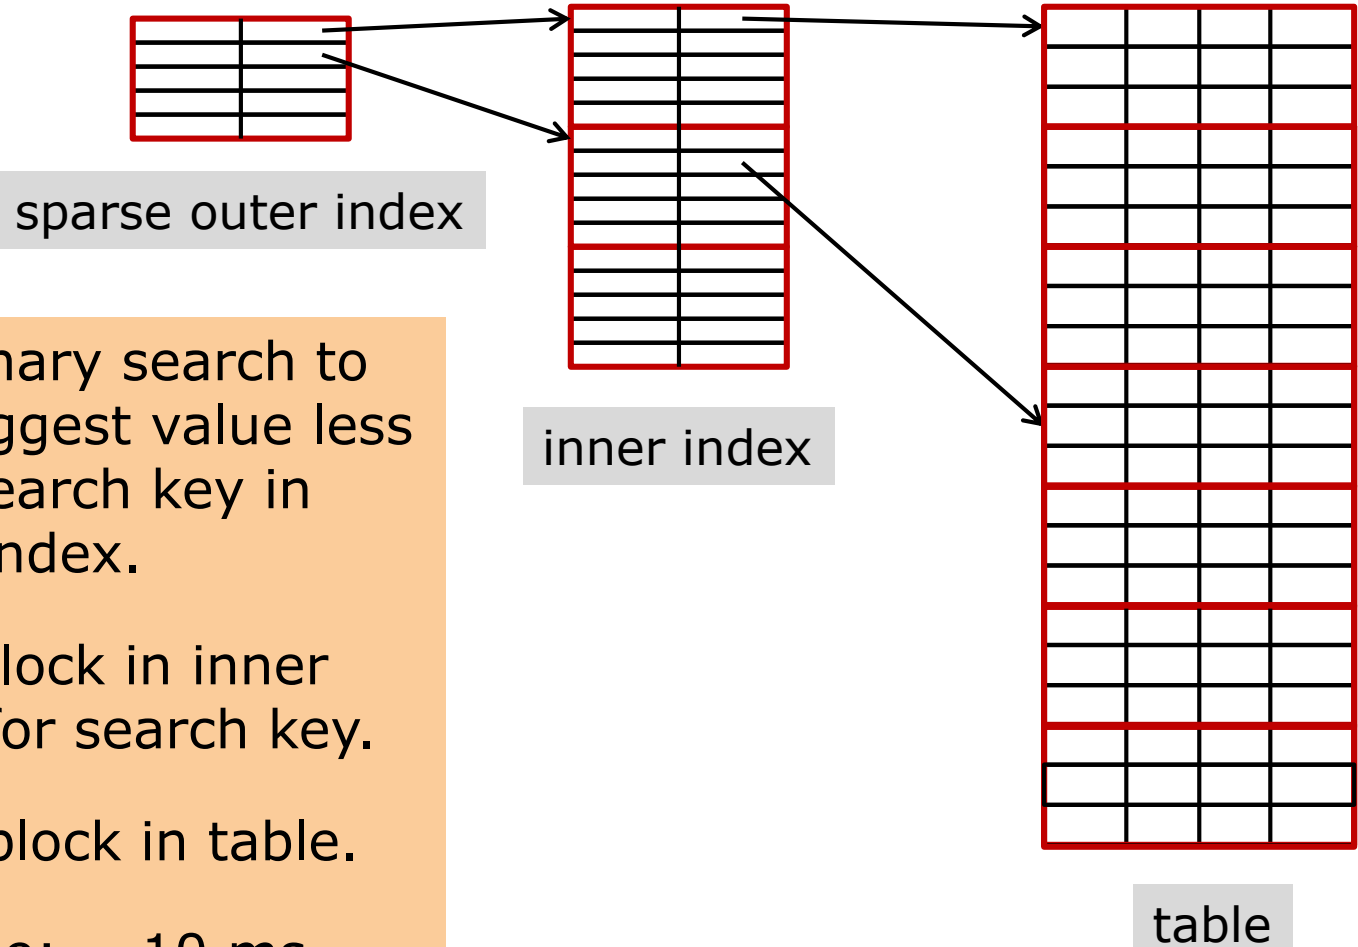
search time: \sim 140 ms



Idea: add an "outer" index




Search process with two indexes



1. Use binary search to find biggest value less than search key in outer index.
2. Scan block in inner index for search key.
3. Go to block in table.

search time: \sim 10 ms

Draw it sideways



Now it's clearly a search tree

The nodes are the size of disk blocks.

If the table gets bigger, more tree levels can be used.

B+ trees

The problem with multi-level indexes:

- As records are inserted and deleted, the index structure needs to be reorganized

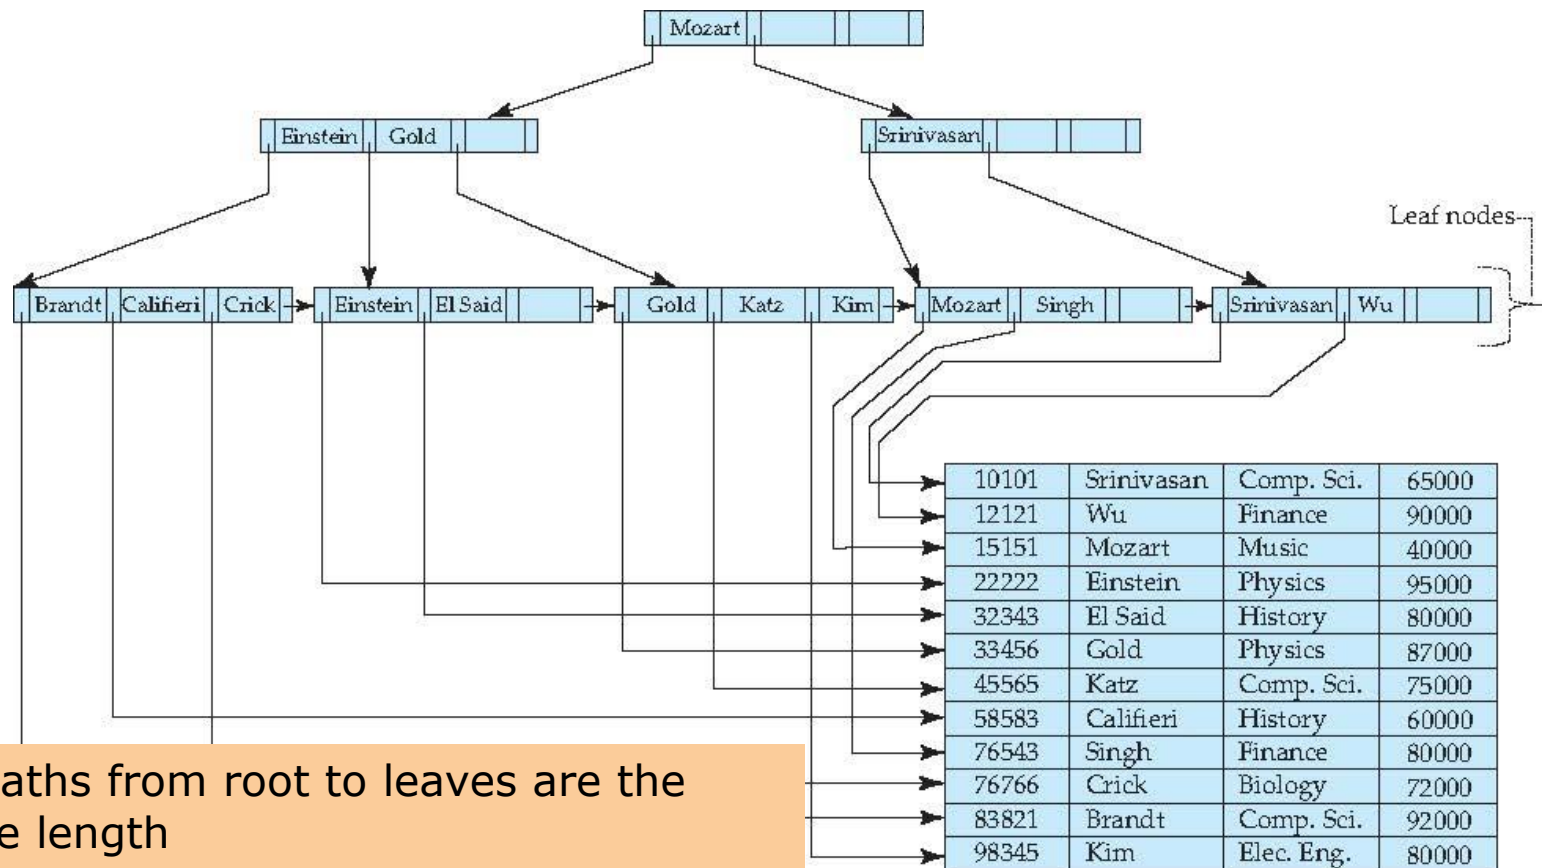
A B+ tree is similar to a multi-level index:

- nodes are the size of disk blocks
- search works from the top to the bottom

but unlike a multi-level index:

- B+ trees are “self-balancing”
- nodes can be partly empty

Features of B+ trees



1. All paths from root to leaves are the same length
2. B+ trees tend to be very "fat" – about 200 children per node

Cost of B+ tree search: roughly the time for 1 block read times the tree depth

(figure from Database System Concepts, Silberschatz, Korth, and Sudarshan)