

# ***DB internals***

---

# First steps

---

We'll assume our system has to process SQL.

We'll assume the data is so large that it can't be stored in memory.

We'll assume the data is stored primarily on a hard drive.

# Query processing

---

Exercise:

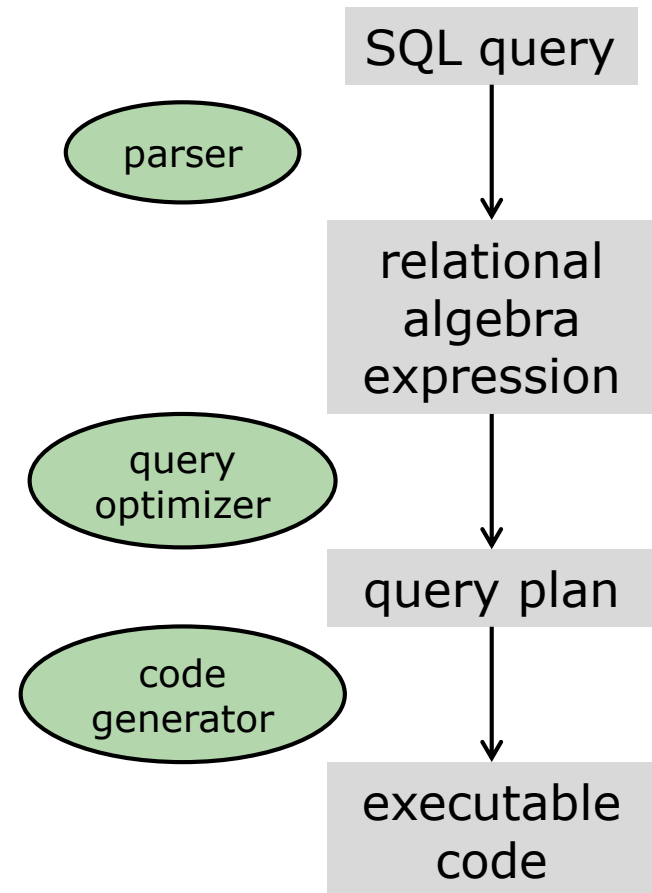
What happens when you run an SQL query?

# Query processing

---

General steps in query processing:

1. parse the SQL
  - check syntax, convert to internal form
2. type check the query
  - are tables, attributes in the database?
3. convert to relational algebra
4. execute the relational algebra



# Query optimization

---

SQL is declarative, so there are many ways to run a query.

- ❑ multiple ways to convert query to relational algebra
- ❑ multiple ways to evaluate a relational algebra expression

# Query optimization

---

A problem is to choose a “query plan”.

How to determine cost of alternative query plans?

Many relevant factors:

- computational time
- I/O time
- required memory/disk space
- impact on time for DB updates

# Estimating the cost of running a query

---

## □ Exercise:

- what is the biggest factor affecting the time it takes to run an SQL query (on a large DB)?

# Estimating the cost of running a query


---

- ❑ Databases are generally too big to put in memory, so we need to use disk (hard drive)
- ❑ Reading from disk is much slower than reading from memory.
- ❑ So, in optimizing query evaluation we want to minimize the number of disk accesses



# Basic of hard drives

- ❑ Data is stored on tracks, sectors
  - 1 track ~ 1000 sectors of 512 bytes each
- ❑ Access time: ~ 10 ms
- ❑ Transfer time: ~ 80 MB/sec
- ❑ Block: unit of transfer between disk and memory
- ❑ Typical block size: 4K or 8K bytes



(figure from <http://www.911-computer.com/how-to-repair-hard-drive-bad-sectors>)

# Reducing disk block access

---

- ❑ store related records in the same block
- ❑ use write buffers

# How to store a database on disk?

---

- To simplify things, we assume:
  - data for each table in its own file
  - all records of a table are the same size
- A **database** = a collection of files
- A **file** is a sequence of blocks
- A **block** is a sequence of records
- A **record** is a sequence of fields

# Example

---

We'll assume  
one file for each  
table.

We can simply  
store the records  
in a file, one  
after another

instructor				
<u>record</u>	ID	name	dept_name	salary
0	10101	Srinivasan	Comp. Sci.	65000
1	12121	Wu	Finance	90000
2	15151	Mozart	Music	40000
3	22222	Einstein	Physics	95000
4	32343	El Said	History	60000
5	33456	Gold	Physics	87000
6	45565	Katz	Comp. Sci.	75000
7	58583	Califieri	History	62000
...	76543	Singh	Finance	80000
	76766	Crick	Biology	72000
	83821	Brandt	Comp. Sci.	92000
	98345	Kim	Elec. Eng.	80000

# Buffering disk data

---

- We want to minimize the number of blocks transferred between disk and memory
- So – blocks are kept in memory if possible
- Buffer manager
  - allocates buffer space (in memory)
  - programs that want data call the buffer manager:
    - if block is in memory, manager returns it
    - otherwise, manager gets block from disk

---

# In-Memory Search

# Idea: linear search

Let's pretend data is in memory.

```
select name from  
student where  
dept_name =  
"Comp. Sci.";
```

Simplest strategy is to search sequentially from top to bottom.

In general we must search **all** records.

**student**

ID	name	dept_name	tot_cred
74732	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
50337	Brandt	History	80
23121	Chavez	Finance	110
35432	Avery	Comp. Sci.	44
44553	Peltier	Physics	56
23145	Levy	Physics	46
54321	Williams	Comp. Sci.	54
88232	Sanchez	Music	38
70557	Snow	Physics	0
20346	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
12543	Bourikas	Elec. Eng.	98
46524	Tanaka	Biology	120




# Special case 1: equality on key

```
select * from  
student where  
ID = 44553;
```

If search is for equality on primary key attribute, stop when found.

Search half the table in average case.

**student**



ID	name	dept_name	tot_cred
74732	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
50337	Brandt	History	80
23121	Chavez	Finance	110
35432	Avery	Comp. Sci.	44
44553	Peltier	Physics	56
23145	Levy	Physics	46
54321	Williams	Comp. Sci.	54
88232	Sanchez	Music	38
70557	Snow	Physics	0
20346	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
12543	Bourikas	Elec. Eng.	98
46524	Tanaka	Biology	120




# Special case 2: ordered field

`select name from  
student where  
total_cred < 50;`

If linear search on  
ordered field, can  
also stop early.

Search half the  
table in average  
case.

**student**



ID	name	dept_name	tot_cred
70557	Snow	Physics	0
12345	Shankar	Comp. Sci.	32
55739	Sanchez	Music	38
35432	Avery	Comp. Sci.	44
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
44553	Peltier	Physics	56
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
19991	Brandt	History	80
98765	Bourikas	Elec. Eng.	98
128	Zhang	Comp. Sci.	102
23121	Chavez	Finance	110
98988	Tanaka	Biology	120

# Linear search: pros and cons

---

## Pros:

- ❑ ultra-simple
- ❑ no special data structures needed
- ❑ no overhead when records added/deleted

## Cons:

- ❑ slow: may need to search entire table

# Binary search


```
select name from  
student where ID  
= 76543;
```

If field is ordered,  
can do **binary  
search**.

## student

ID	name	dept_name	tot_cred
128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
35432	Avery	Comp. Sci.	44
44553	Peltier	Physics	56
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
55739	Sanchez	Music	38
70557	Snow	Physics	0
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
98765	Bourikas	Elec. Eng.	98
98988	Tanaka	Biology	120

# Binary search



## student

ID	name	dept_name	tot_cred
128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
35432	Avery	Comp. Sci.	44
44553	Peltier	Physics	56
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
55739	Sanchez	Music	38
70557	Snow	Physics	0
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
98765	Bourikas	Elec. Eng.	98
98988	Tanaka	Biology	120

# Binary search: pros and cons

---

## Pros:

- ❑ pretty simple
- ❑ much faster than linear search

## Cons:

- ❑ only one field of a table can be sorted
- ❑ overhead to keep data sorted when rows are added, deleted, modified

---

# On-Disk Search

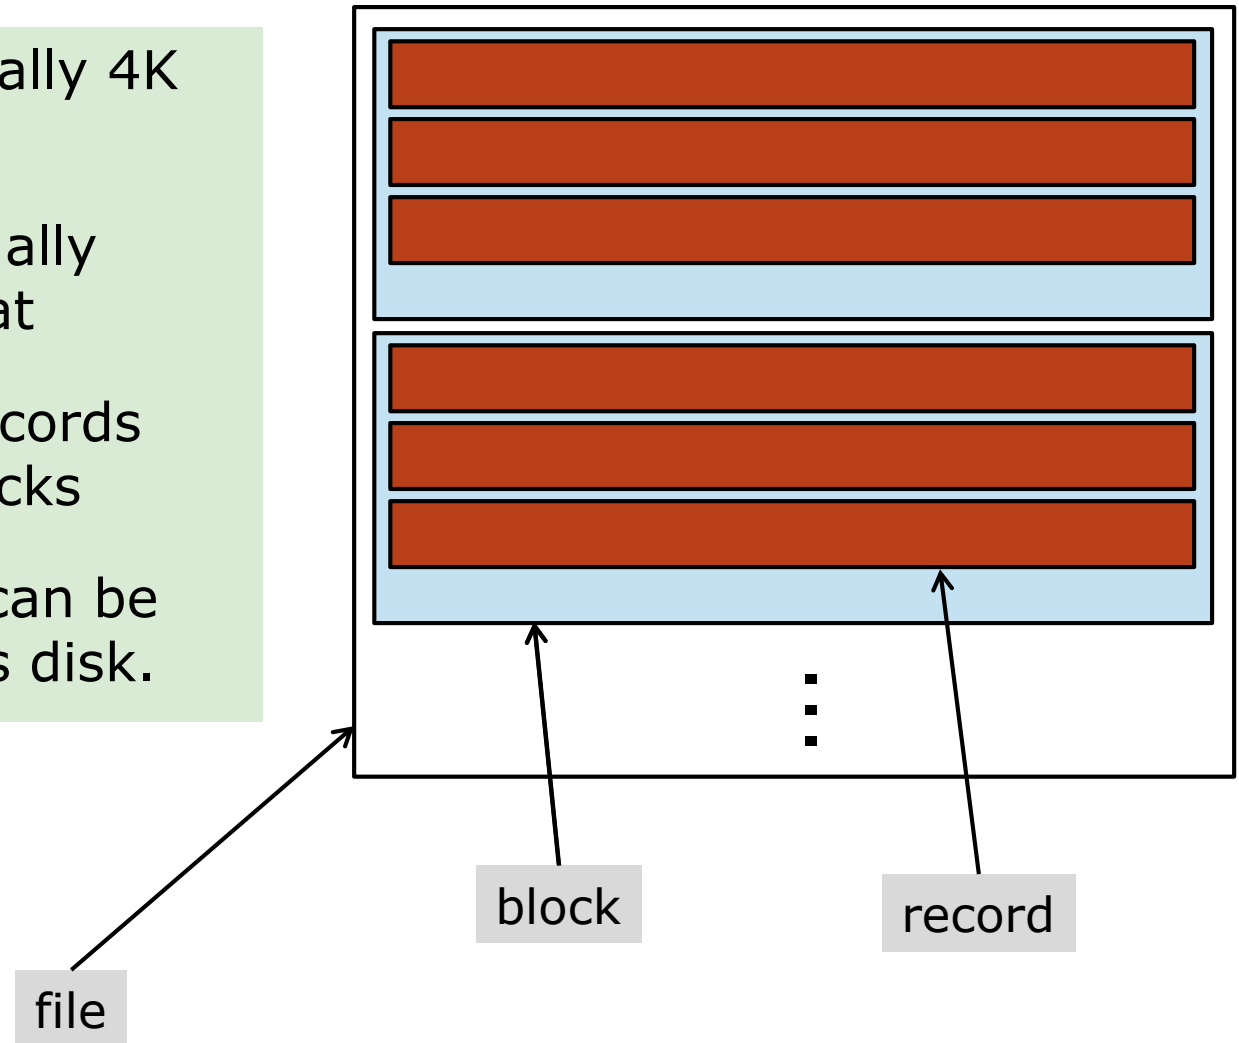
# Records and blocks

Blocks are typically 4K or 8K bytes

Records are usually smaller than that

We'll assume records don't bridge blocks

Blocks of a file can be scattered across disk.



# Linear search over table on disk

---

With data on disk, our cost estimate is based on time spent on disk I/O.

Linear search, *assuming blocks are stored contiguously on disk*:

$$\text{access\_time} + \text{num\_blocks} * \text{transfer\_time}$$

typical access\_time: 4 ms

typical transfer\_time: 0.1 ms

typical block size: 4K or 8K bytes

Note: transfer\_time is the time to transfer one block.  
It can be derived from the disk transfer if you know the block size.



# Binary search

---

With binary search, we will have to “skip around” the disk, even if file blocks are contiguous.

Also, without additional data structures, how do you find the address of a block in the “middle” of some rows?

# File organization

---

How are the records ordered within a file?

**heap** organization:

- records are in an arbitrary order

**sequential** organization

- records are ordered according to the value of one attribute
- this attribute is called the “search key”

# Sequential files

Ideally, in sequential organization, records are physically ordered according to the search key.

How to cope with insertion and deletion of records?


## student

ID	name	dept_name	tot_cred
128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
35432	Avery	Comp. Sci.	44
44553	Peltier	Physics	56
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
55739	Sanchez	Music	38
70557	Snow	Physics	0
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
98765	Bourikas	Elec. Eng.	98
98988	Tanaka	Biology	120

# Sequential file organization

## instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000



A pointer in each record points to the next record.

Essentially a linked list.

Every so often the file is **reorganized**.