

NoSQL : Databases at web scale

Relational DBs solve real problems

Relational DBs are a very good solution to these problems:

- ❑ How to support concurrent access to the data by thousands of users?
- ❑ How to keep data consistent, and hide temporary inconsistency?
- ❑ How to avoid redundancy?
- ❑ How to let multiple applications share data?

New requirements

Trends since the '90s have led to new requirements:

- ❑ Need to support millions or 10s of millions of users
- ❑ Need to support rapid feature development and schema changes
- ❑ Need to access data from OO languages
- ❑ Need for very high availability

Can relational DBs scale?

- ❑ vertical scaling (with relational DBs)
 - get bigger, faster machines, special hardware
 - very expensive; may not provide enough scaling
- ❑ horizontal scaling
 - distribute data over many cheap machines
 - sharding

Schema changes

- ❑ Schema changes are painful with relational DBs
- ❑ “Data migration”
- ❑ Data migration might involve:
 - moving data off database
 - modifying schema
 - modifying data
 - putting data back in database
- ❑ Problem is worse when database shared by many applications

Availability and Consistency

Many applications want “availability, at the cost of consistency”

Examples:

- Amazon cart

When databases get distributed, and failures can occur, you can't have both availability and consistency.

Cost factor

Commercial databases are expensive.

- proof: Larry Ellison (Oracle) is worth \$52 billion

The main NoSQL databases are open source.

What is the alternative?

- ❑ commodity hardware
- ❑ large clusters
- ❑ open source
- ❑ no schema
- ❑ simple API
- ❑ eventual consistency



NoSQL

NoSQL databases

There are different kinds of NoSQL databases.

Main types:

type	examples
key-value	Redis, DynamoDB
document	MongoDB, Couchbase
column-family	Cassandra, HBase
graph	Neo4J, OrientDB

Key-value

- ❑ super simple API
- ❑ if you want to store something, give it a name and store it
 - database does not “know” the structure of the data
- ❑ core data structure is the associative array
- ❑ core technology is distributed hash table

In an associative array:

- different types of values can be stored
- any type of value can be used as an index

Document

- ❑ Similar to key-value, but now values have standard structure
- ❑ Typically values are XML or JSON documents
- ❑ Closer to relational DB; concepts like “collections”
- ❑ Supports a query language
- ❑ No pre-defined schema

Column-family

- ❑ Another variant of key-value store
- ❑ Data actually stored in columns, not rows
 - or “families of columns”
- ❑ Reads and writes are done on columns, not rows
- ❑ Rows can have different number of columns
- ❑ Search through a column becomes very fast; high performance for many queries

Graph

- ❑ Data is modeled as a graph
- ❑ Key-value store extended to allow relationships between the values
- ❑ Data stored in a key-value store
- ❑ Supports efficient traversal between values




figure from: neo4j.com/developer/graph-db-vs-nosql/

When to use NoSQL

Factors in favor of NoSQL:

- data is so big that a cluster is required
- data is non-uniform
- data is in the form of aggregates
- availability is important

Factors in favor of relational DBs:

- consistency is important
- a large and stable set of tools are needed
- many applications will use the data
- security is important

Summary

NoSQL arose because of new requirements:

- millions of users
- high availability
- rapid app development

Main types of NoSQL databases:

- key-value
- document
- column family
- graph